

Statistics

Technical reference

Companion to the statistics notes on lorenzoliuzzo.github.io. Probability, standard random variables, sampling distributions, estimation, intervals and tests, with proofs.

Contents

1	Probability toolkit	2
1.1	Probability spaces and the axioms	2
1.2	Conditional probability, independence, Bayes	2
1.3	Random variables and distributions	3
1.4	Expectation	3
1.5	Variance, covariance, correlation	4
1.6	Markov and Chebyshev	4
1.7	Limit theorems	4
2	Discrete random variables	5
2.1	Bernoulli and binomial	5
2.2	Geometric	5
2.3	Poisson	6
3	Continuous random variables	6
3.1	Uniform and inverse transform sampling	6
3.2	Exponential	6
3.3	Normal	7
3.4	Gamma and chi-square	7
4	Sampling distributions of statistics	7
4.1	Mean, variance and proportion	7
4.2	Normal-theory distributions	8
4.3	The t and F distributions in more detail	8
4.4	A non-normal example: the sample maximum	9
4.5	Simulation and the bootstrap	9
5	Point estimation	9
5.1	Fisher information and the Cramér–Rao bound	10
5.2	Sufficiency (brief)	10
6	Maximum likelihood	10
6.1	Asymptotic properties	11
6.2	Method of moments	11
7	Interval estimation	11
8	Hypothesis testing	12
8.1	Framework	12
8.2	The t -tests	12
8.3	Power and sample size	13
8.4	Likelihood ratio tests	13
8.5	Multiple testing	13
9	Summary of the main formulas	13
10	References	14

1 Probability toolkit

1.1 Probability spaces and the axioms

A probability space is a triple $(\Omega, \mathcal{F}, P)$: a sample space Ω , a σ -algebra $\mathcal{F}$ of events (a family of subsets of Ω containing Ω and closed under complements and countable unions), and a function $P : \mathcal{F} \rightarrow [0, 1]$ satisfying the axioms below. Events A, B are *disjoint* if $A \cap B = \emptyset$; events $B_1, B_2, \dots$ form a *partition* if they are pairwise disjoint and $\bigcup_i B_i = \Omega$.

Definition (Kolmogorov axioms). (A1) $P(A) \geq 0$ for every event A . (A2) $P(\Omega) = 1$. (A3) For pairwise disjoint events $A_1, A_2, \dots$, $P(\bigcup_i A_i) = \sum_i P(A_i)$.

Theorem (Consequences of the axioms). For events $A, B, A_1, A_2, \dots$: (i) $P(\emptyset) = 0$; (ii) finite additivity: $P(\bigcup_{i=1}^n A_i) = \sum_{i=1}^n P(A_i)$ for disjoint A_i ; (iii) $P(A^c) = 1 - P(A)$; (iv) if $A \subseteq B$ then $P(A) \leq P(B)$; in particular $P(A) \leq 1$; (v) $P(A \cup B) = P(A) + P(B) - P(A \cap B)$; (vi) union bound: $P(\bigcup_i A_i) \leq \sum_i P(A_i)$.

Proof. (i) Take $A_i = \emptyset$ for all i in (A3): $P(\emptyset) = \sum_i P(\emptyset)$, which forces $P(\emptyset) = 0$. (ii) Apply (A3) with $A_i = \emptyset$ for $i > n$ and use (i). (iii) $\Omega = A \cup A^c$ is a disjoint union, so $1 = P(A) + P(A^c)$. (iv) $B = A \cup (B \setminus A)$ is disjoint, so $P(B) = P(A) + P(B \setminus A) \geq P(A)$; with $B = \Omega$ this gives $P(A) \leq 1$. (v) Write $A \cup B = A \cup (B \setminus A)$ and $B = (A \cap B) \cup (B \setminus A)$, both disjoint; then $P(B \setminus A) = P(B) - P(A \cap B)$, and substituting gives (v). (vi) Set $B_1 = A_1$ and $B_i = A_i \setminus \bigcup_{j < i} A_j$. The B_i are disjoint, $\bigcup_i B_i = \bigcup_i A_i$ and $B_i \subseteq A_i$, so by (A3) and (iv) $P(\bigcup_i A_i) = \sum_i P(B_i) \leq \sum_i P(A_i)$. $\square$

The general inclusion–exclusion formula is $P(\bigcup_{i=1}^n A_i) = \sum_i P(A_i) - \sum_{i < j} P(A_i \cap A_j) + \sum_{i < j < k} P(A_i \cap A_j \cap A_k) - \dots$. Two further facts follow from (A3) by taking limits: if A_n increases to A then $P(A_n) \rightarrow P(A)$, and likewise for decreasing sequences (continuity of P). When Ω is finite and outcomes are equally likely, $P(A) = |A| / |\Omega|$.

1.2 Conditional probability, independence, Bayes

Definition. For $P(B) > 0$, the conditional probability of A given B is $P(A | B) = P(A \cap B) / P(B)$. Then $P(\cdot | B)$ is itself a probability: it satisfies (A1)–(A3) on $\mathcal{F}$, so all the consequences above hold conditionally.

Rearranging gives the multiplication rule $P(A \cap B) = P(A | B) P(B)$ and, by induction, the chain rule

$$P(A_1 \cap \dots \cap A_n) = P(A_1) P(A_2 | A_1) \dots P(A_n | A_1 \cap \dots \cap A_{n-1}). \quad (1)$$

Definition (Independence). A and B are independent if $P(A \cap B) = P(A) P(B)$ (equivalently $P(A | B) = P(A)$ when $P(B) > 0$). Events $A_1, \dots, A_n$ are *mutually independent* if $P(\bigcap_{i \in S} A_i) = \prod_{i \in S} P(A_i)$ for every subset S of indices; pairwise independence alone does not imply this. A and B are conditionally independent given C if $P(A \cap B | C) = P(A | C) P(B | C)$.

Remark. Disjoint events with positive probability are never independent: $P(A \cap B) = 0 \neq P(A) P(B)$.

Theorem (Law of total probability and Bayes' theorem). Let $B_1, B_2, \dots$ be a partition of Ω with $P(B_i) > 0$. For any event A ,

$$P(A) = \sum_i P(A | B_i) P(B_i), \quad (2)$$

and, if $P(A) > 0$,

$$P(B_j | A) = \frac{P(A | B_j) P(B_j)}{\sum_i P(A | B_i) P(B_i)}. \quad (3)$$

Proof. The sets $A \cap B_i$ are disjoint with union A , so by (A3) and the multiplication rule $P(A) = \sum_i P(A \cap B_i) = \sum_i P(A | B_i) P(B_i)$. For Bayes, $P(B_j | A) = P(A \cap B_j) / P(A) = P(A | B_j) P(B_j) / P(A)$, and the denominator is Equation 2. $\square$

In Bayesian language $P(B_j)$ is the prior, $P(A | B_j)$ the likelihood, $P(A)$ the evidence (marginal likelihood) and $P(B_j | A)$ the posterior. The odds form is $P(B | A) / P(B^c | A) = [P(A | B) / P(A | B^c)] \cdot [P(B) / P(B^c)]$: posterior odds equal likelihood ratio times prior odds.

Example (Diagnostic test). Prevalence $P(D) = 0.01$, sensitivity $P(+ | D) = 0.95$, false-positive rate $P(+ | D^c) = 0.10$. Then

$$P(D | +) = \frac{0.95 \cdot 0.01}{0.95 \cdot 0.01 + 0.10 \cdot 0.99} = \frac{0.0095}{0.1085} \approx 0.088. \quad (4)$$

The posterior is below 9% because the healthy majority produces more false positives than the sick minority produces true ones: the base rate matters.

1.3 Random variables and distributions

A random variable is a measurable function $X : \Omega \rightarrow \mathbb{R}$. Its distribution function is $F(x) = P(X \leq x)$.

Lemma (Properties of the distribution function). F is non-decreasing and right-continuous, with $\lim_{x \rightarrow -\infty} F(x) = 0$ and $\lim_{x \rightarrow \infty} F(x) = 1$. For $a < b$, $P(a < X \leq b) = F(b) - F(a)$.

Proof. Monotonicity is (iv) applied to $\{X \leq x\} \subseteq \{X \leq y\}$. Right-continuity and the limits follow from continuity of P along monotone sequences of events. Finally $\{X \leq b\} = \{X \leq a\} \cup \{a < X \leq b\}$ is disjoint. $\square$

X is discrete with mass function $p(x) = P(X = x)$, $\sum_x p(x) = 1$, or continuous with density $f = F'$, $\int f = 1$ and $P(a \leq X \leq b) = \int_a^b f$; in the continuous case $P(X = x) = 0$ for every x . The quantile function is $F^{-1}(p) = \inf\{x : F(x) \geq p\}$ and the median is $F^{-1}(1/2)$. For a pair (X, Y) with joint density $f_{X,Y}$, the marginal is $f_{X(x)} = \int f_{X,Y}(x, y) dy$, the conditional density is $f_{X|Y}(x | y) = f_{X,Y}(x, y) / f_{Y(y)}$, and X, Y are independent iff $f_{X,Y} = f_X f_Y$.

1.4 Expectation

Definition. $E[X] = \sum_x x p(x)$ (discrete) or $\int x f(x) dx$ (continuous), when the sum or integral converges absolutely. For a function g (law of the unconscious statistician), $E[g(X)] = \sum_x g(x) p(x)$ or $\int g(x) f(x) dx$.

Theorem (Properties of expectation). (i) Linearity: $E[aX + bY + c] = a E[X] + b E[Y] + c$ for any X, Y , dependent or not. (ii) Monotonicity: $X \leq Y$ implies $E[X] \leq E[Y]$. (iii) If X, Y are independent, $E[XY] = E[X] E[Y]$. (iv) $E[\mathbb{1}_A] = P(A)$. (v) Jensen: for convex g , $g(E[X]) \leq E[g(X)]$.

Proof. (i) and (ii) follow from linearity and monotonicity of sums and integrals, using the joint density for X and Y together. (iii) By independence $f_{X,Y} = f_X f_Y$, so $E[XY] = \int \int xy f_{X(x)} f_{Y(y)} dx dy = E[X] E[Y]$. (iv) $\mathbb{1}_A$ takes the value 1 with probability $P(A)$ and 0 otherwise. (v) A convex g lies above its tangent line at $m = E[X]$: $g(x) \geq g(m) + c(x - m)$ for some slope c ; take expectations. $\square$

1.5 Variance, covariance, correlation

Definition. $\text{Var}(X) = \mathbb{E}[(X - \mu)^2]$ with $\mu = \mathbb{E}[X]$; the standard deviation is $\sigma = \sqrt{\text{Var}(X)}$. $\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)]$ and, when $\sigma_X, \sigma_Y > 0$, the correlation is $\rho_{XY} = \text{Cov}(X, Y)/(\sigma_X \sigma_Y)$. Moments: $\mathbb{E}[X^k]$; central moments $\mathbb{E}[(X - \mu)^k]$; skewness $\mathbb{E}[(X - \mu)^3]/\sigma^3$ and kurtosis $\mathbb{E}[(X - \mu)^4]/\sigma^4$ (equal to 3 for the normal).

Theorem (Properties of variance and covariance). (i) $\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}X)^2 \geq 0$, with equality iff X is almost surely constant. (ii) $\text{Var}(aX + b) = a^2 \text{Var}(X)$. (iii) $\text{Cov}(X, Y) = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y]$, $\text{Cov}(X, X) = \text{Var}(X)$, and Cov is bilinear. (iv) $\text{Var}(aX + bY) = a^2 \text{Var}(X) + b^2 \text{Var}(Y) + 2ab \text{Cov}(X, Y)$. (v) Independent X, Y have $\text{Cov}(X, Y) = 0$; the converse is false. (vi) $|\rho_{XY}| \leq 1$, with equality iff $Y = aX + b$ almost surely.

Proof. (i) Expand $(X - \mu)^2 = X^2 - 2\mu X + \mu^2$ and take expectations; nonnegativity holds because the integrand is ≥ 0 , and $\mathbb{E}[(X - \mu)^2] = 0$ forces $X = \mu$ almost surely. (ii) $aX + b - \mathbb{E}[aX + b] = a(X - \mu)$. (iii) Expand the product and use linearity; bilinearity is linearity in each argument. (iv) is (ii) and (iii) combined. (v) follows from (iii) and property (iii) of expectation. For the failure of the converse, let $X \sim \mathcal{N}(0, 1)$ and $Y = X^2$: $\text{Cov}(X, Y) = \mathbb{E}[X^3] - \mathbb{E}[X]\mathbb{E}[X^2] = 0$ since odd moments of a symmetric law vanish, yet Y is a function of X . (vi) For every real t , $0 \leq \text{Var}(Y - tX) = \text{Var}(Y) - 2t \text{Cov}(X, Y) + t^2 \text{Var}(X)$. A quadratic in t that never goes negative has non-positive discriminant, $4\text{Cov}(X, Y)^2 - 4 \text{Var}(X) \text{Var}(Y) \leq 0$, which is $|\rho| \leq 1$. Equality means the discriminant vanishes, so $\text{Var}(Y - tX) = 0$ for some t and $Y - tX$ is constant. $\square$

Property (iv) generalizes to $\text{Var}(\sum_i X_i) = \sum_i \text{Var}(X_i) + 2 \sum_{i < j} \text{Cov}(X_i, X_j)$, and so, for uncorrelated (hence independent) X_i with common variance σ^2 ,

$$\text{Var}\left(\sum_i X_i\right) = n\sigma^2, \quad \text{Var}(\bar{X}_n) = \sigma^2/n. \quad (5)$$

Theorem (Law of total expectation and total variance). If $\mathbb{E}[X^2] < \infty$, then $\mathbb{E}[X] = \mathbb{E}[\mathbb{E}[X | Y]]$ and

$$\text{Var}(X) = \mathbb{E}[\text{Var}(X | Y)] + \text{Var}(\mathbb{E}[X | Y]). \quad (6)$$

Proof. The first identity is the law of total probability for expectations: $\mathbb{E}[X] = \int \mathbb{E}[X | Y = y] f_{Y(y)} dy$. For the second, write $X - \mathbb{E}X = (X - \mathbb{E}[X | Y]) + (\mathbb{E}[X | Y] - \mathbb{E}X)$ and square. The cross term has expectation zero: conditionally on Y the second factor is fixed and $\mathbb{E}[X - \mathbb{E}[X | Y] | Y] = 0$. The first square has expectation $\mathbb{E}[\text{Var}(X | Y)]$ and the second $\text{Var}(\mathbb{E}[X | Y])$. $\square$

1.6 Markov and Chebyshev

Lemma (Markov and Chebyshev). If $X \geq 0$ and $a > 0$ then $\mathbb{P}(X \geq a) \leq \mathbb{E}[X]/a$. For any X with mean μ and variance σ^2 and any $k > 0$, $\mathbb{P}(|X - \mu| \geq k) \leq \sigma^2/k^2$.

Proof. $a \mathbb{1}\{X \geq a\} \leq X$ pointwise, take expectations. Chebyshev is Markov applied to $(X - \mu)^2$ and $a = k^2$. $\square$

The standard discrete and continuous distributions, with their means, variances and uses, are in the chapters on discrete and continuous random variables below.

1.7 Limit theorems

Write $\bar{X}_n = n^{-1} \sum_{i=1}^n X_i$.

Theorem (Weak law of large numbers). If X_i are i.i.d. with $E X_1 = \mu$ and $\text{Var}(X_1) = \sigma^2 < \infty$, then $\bar{X}_n \xrightarrow{P} \mu$, i.e. $P(|\bar{X}_n - \mu| > \varepsilon) \rightarrow 0$ for every $\varepsilon > 0$.

Proof. By Equation 5, $\text{Var}(\bar{X}_n) = \sigma^2/n$. Chebyshev gives $P(|\bar{X}_n - \mu| \geq \varepsilon) \leq \sigma^2/(n\varepsilon^2) \rightarrow 0$. $\square$

Theorem (Central limit theorem). If X_i are i.i.d. with mean μ and variance $0 < \sigma^2 < \infty$, then

$$\sqrt{n}(\bar{X}_n - \mu)/\sigma \xrightarrow{d} \mathcal{N}(0, 1). \quad (7)$$

The proof goes through characteristic functions and is not reproduced here. The Berry–Esseen theorem quantifies the rate: the Kolmogorov distance to Φ is at most $C E|X_1 - \mu|^3 / (\sigma^3 \sqrt{n})$. This is why the sample size needed depends on skewness.

Theorem (Slutsky and delta method). If $A_n \xrightarrow{d} A$ and $B_n \xrightarrow{P} b$ (constant), then $A_n + B_n \xrightarrow{d} A + b$ and $A_n B_n \xrightarrow{d} bA$. If moreover $\sqrt{n}(T_n - \theta) \xrightarrow{d} \mathcal{N}(0, \tau^2)$ and g is differentiable at θ with $g'(\theta) \neq 0$, then

$$\sqrt{n}(g(T_n) - g(\theta)) \xrightarrow{d} \mathcal{N}(0, g'(\theta)^2 \tau^2). \quad (8)$$

Slutsky is what licenses replacing an unknown σ by a consistent estimate in a pivot, which is how the Wald and large-sample t intervals are justified.

2 Discrete random variables

A discrete random variable has mass function $p(x) = P(X = x)$ on a countable support, with $\sum_x p(x) = 1$, $E[X] = \sum_x x p(x)$ and $\text{Var}(X) = E[X^2] - (E[X])^2$. Probability background is in the Probability toolkit chapter.

2.1 Bernoulli and binomial

Definition. $X \sim \text{Bernoulli}(p)$ if $P(X = 1) = p$ and $P(X = 0) = 1 - p$. $Y \sim \text{Binomial}(n, p)$ is the sum of n independent $\text{Bernoulli}(p)$ variables, with mass function $P(Y = k) = \binom{n}{k} p^k (1 - p)^{n-k}$, $k = 0, \dots, n$.

Theorem (Mean and variance). If $X \sim \text{Bernoulli}(p)$ then $E[X] = p$ and $\text{Var}(X) = p(1 - p)$. If $Y \sim \text{Binomial}(n, p)$ then $E[Y] = np$ and $\text{Var}(Y) = np(1 - p)$.

Proof. $E[X] = 1 \cdot p + 0 \cdot (1 - p) = p$, and since $X^2 = X$, $E[X^2] = p$ and $\text{Var}(X) = p - p^2$. For $Y = \sum_{i=1}^n X_i$ with independent X_i , linearity gives $E[Y] = np$ and independence gives $\text{Var}(Y) = \sum_i \text{Var}(X_i) = np(1 - p)$. $\square$

The binomial mass function counts the $\binom{n}{k}$ arrangements of k successes among n trials, each with probability $p^k (1 - p)^{n-k}$. For large n it is approximately $\mathcal{N}(np, np(1 - p))$ by the central limit theorem; a common rule is $np(1 - p) \geq 10$.

2.2 Geometric

Definition. $X \sim \text{Geometric}(p)$ counts the trials up to and including the first success: $P(X = k) = (1 - p)^{k-1} p$, $k = 1, 2, \dots$

Theorem (Mean and memorylessness). $E[X] = 1/p$ and $\text{Var}(X) = (1 - p)/p^2$. Moreover $P(X > m + k \mid X > m) = P(X > k)$ for all integers $m, k \geq 0$.

Proof. With $q = 1 - p$, $E[X] = p \sum_{k \geq 1} k q^{k-1} = p/(1 - q)^2 = 1/p$, using $\sum_{k \geq 1} k q^{k-1} = \frac{d}{dq} \sum_{k \geq 0} q^k = (1 - q)^{-2}$. A second differentiation gives $E[X(X - 1)]$ and hence the variance. For the tail, $P(X > k) = q^k$, so $P(X > m + k \mid X > m) = q^{m+k}/q^m = q^k$. $\square$

2.3 Poisson

Definition. $X \sim \text{Poisson}(\lambda)$ has $P(X = k) = e^{-\lambda} \lambda^k / k!$, $k = 0, 1, 2, \dots$

Theorem (Moments, sums and the binomial limit). (i) $E[X] = \text{Var}(X) = \lambda$. (ii) If $X \sim \text{Poisson}(\lambda)$ and $Y \sim \text{Poisson}(\mu)$ are independent then $X + Y \sim \text{Poisson}(\lambda + \mu)$. (iii) If $X_n \sim \text{Binomial}(n, \lambda/n)$ then $P(X_n = k) \rightarrow e^{-\lambda} \lambda^k / k!$ for every fixed k .

Proof. (i) $E[X] = \sum_{k \geq 1} k e^{-\lambda} \lambda^k / k! = \lambda e^{-\lambda} \sum_{k \geq 1} \lambda^{k-1} / (k-1)! = \lambda$, and similarly $E[X(X - 1)] = \lambda^2$, so $\text{Var}(X) = \lambda^2 + \lambda - \lambda^2 = \lambda$. (ii) $P(X + Y = m) = \sum_{j=0}^m e^{-\lambda} \lambda^j / j! \cdot e^{-\mu} \mu^{m-j} / (m - j)! = e^{-(\lambda + \mu)} (\lambda + \mu)^m / m!$ by the binomial theorem. (iii) $P(X_n = k) = [n! / ((n - k)! k!)] \cdot (\lambda/n)^k (1 - \lambda/n)^{n-k}$. As $n \rightarrow \infty$ the first factor tends to $1/k!$, $(1 - \lambda/n)^n \rightarrow e^{-\lambda}$ and $(1 - \lambda/n)^{-k} \rightarrow 1$. $\square$

3 Continuous random variables

A continuous random variable has a density f with $P(a \leq X \leq b) = \int_a^b f$, mean $E[X] = \int x f(x) dx$ and distribution function $F(x) = \int_{-\infty}^x f$.

3.1 Uniform and inverse transform sampling

$X \sim \text{Uniform}(a, b)$ has $f(x) = 1/(b - a)$ on $[a, b]$, $E[X] = (a + b)/2$ and $\text{Var}(X) = (b - a)^2/12$ (compute $E[X^2] = (a^2 + ab + b^2)/3$).

Theorem (Inverse transform). Let F be a distribution function, $F^{-1}(u) = \inf\{x : F(x) \geq u\}$, and $U \sim \text{Uniform}(0, 1)$. Then $F^{-1}(U)$ has distribution function F . Conversely, if X has a continuous distribution function F , then $F(X) \sim \text{Uniform}(0, 1)$.

Proof. The generalized inverse satisfies $F^{-1}(u) \leq x$ iff $u \leq F(x)$, so $P(F^{-1}(U) \leq x) = P(U \leq F(x)) = F(x)$. For the converse, with F continuous and strictly increasing on the support, $P(F(X) \leq u) = P(X \leq F^{-1}(u)) = F(F^{-1}(u)) = u$. $\square$

The converse is the probability integral transform used for the uniformity of p -values; the direct statement is how simulation software turns uniform random numbers into draws from any distribution with a computable F^{-1} .

3.2 Exponential

Definition. $X \sim \text{Exponential}(\lambda)$ has $f(x) = \lambda e^{-\lambda x}$, $x \geq 0$, so that $P(X > x) = e^{-\lambda x}$.

Theorem (Moments and memorylessness). $E[X] = 1/\lambda$, $\text{Var}(X) = 1/\lambda^2$, and $P(X > s + t \mid X > s) = P(X > t)$.

Proof. Integrating by parts, $E[X] = \int_0^\infty x \lambda e^{-\lambda x} dx = 1/\lambda$ and $E[X^2] = 2/\lambda^2$, so $\text{Var}(X) = 1/\lambda^2$. For the tail, $P(X > s + t \mid X > s) = e^{-\lambda(s+t)}/e^{-\lambda s} = e^{-\lambda t}$. $\square$

The exponential is the only continuous distribution on $[0, \infty)$ with this property. If events arrive as a Poisson process of rate λ (the count in an interval of length t is $\text{Poisson}(\lambda t)$), the waiting time to the first event, and between successive events, is $\text{Exponential}(\lambda)$: $P(T > t) = P(\text{no event in } [0, t]) = e^{-\lambda t}$.

3.3 Normal

Definition. $X \sim \mathcal{N}(\mu, \sigma^2)$ has density $(2\pi\sigma^2)^{-1/2} \exp(-(x - \mu)^2/(2\sigma^2))$. The standard normal $\mathcal{N}(0, 1)$ has distribution function Φ and quantiles $z_p = \Phi^{-1}(p)$.

Theorem (Standardization and closure). (i) $E[X] = \mu$ and $\text{Var}(X) = \sigma^2$. (ii) $Z = (X - \mu)/\sigma \sim \mathcal{N}(0, 1)$, so $P(X \leq x) = \Phi((x - \mu)/\sigma)$. (iii) If $X_1, \dots, X_n$ are independent, $X_i \sim \mathcal{N}(\mu_i, \sigma_i^2)$, then $\sum_i a_i X_i + b \sim \mathcal{N}(\sum_i a_i \mu_i + b, \sum_i a_i^2 \sigma_i^2)$.

Proof. (ii) The substitution $y = (u - \mu)/\sigma$ in $\int_{-\infty}^x f(u) du$ gives $\int_{-\infty}^{(x-\mu)/\sigma} \varphi(y) dy$. (i) follows from (ii) since $E[Z] = 0$ (odd integrand) and $E[Z^2] = 1$ (integration by parts). (iii) The mean and variance follow from linearity and independence; that the sum is again normal follows from the moment generating function $E[e^{tX}] = e^{\mu t + \sigma^2 t^2/2}$, which multiplies for independent summands. $\square$

Quantiles that recur: $z_{0.90} \approx 1.282$, $z_{0.95} \approx 1.645$, $z_{0.975} \approx 1.960$, $z_{0.995} \approx 2.576$, and $P(|Z| \leq 1, 2, 3) \approx 0.683, 0.954, 0.997$.

3.4 Gamma and chi-square

Definition. $X \sim \text{Gamma}(\alpha, \beta)$ (shape $\alpha > 0$, rate $\beta > 0$) has density $\beta^\alpha x^{\alpha-1} e^{-\beta x} / \Gamma(\alpha)$ on $x > 0$, with $E[X] = \alpha/\beta$ and $\text{Var}(X) = \alpha/\beta^2$. The sum of k independent $\text{Exponential}(\lambda)$ variables is $\text{Gamma}(k, \lambda)$. The chi-square distribution is $\chi_k^2 = \text{Gamma}(k/2, 1/2)$, with density $x^{k/2-1} e^{-x/2} / (2^{k/2} \Gamma(k/2))$.

Theorem (Chi-square as a sum of squares). If $Z_1, \dots, Z_k$ are independent $\mathcal{N}(0, 1)$ then $\sum_i Z_i^2 \sim \chi_k^2$, with mean k and variance $2k$.

Proof. The density of Z_1^2 follows by the change of variable $y = z^2$ and equals the χ_1^2 density; independent Gamma variables with a common rate add their shapes, giving χ_k^2 . For the moments, $E[Z^2] = 1$ and $\text{Var}(Z^2) = E[Z^4] - 1 = 3 - 1 = 2$, and the Z_i^2 are independent, so the mean is k and the variance $2k$. $\square$

4 Sampling distributions of statistics

A statistic $T = T(X_1, \dots, X_n)$ is a random variable. Its sampling distribution is the law of T when the sample is drawn from the assumed model, and its standard error is $\text{se}(T) = \sqrt{\text{Var}(T)}$.

4.1 Mean, variance and proportion

Theorem (Mean and variance of $\bar{X}$). For i.i.d. X_i with mean μ and variance σ^2 , $E[\bar{X}] = \mu$ and $\text{Var}(\bar{X}) = \sigma^2/n$. For two independent samples, $\bar{X}_1 - \bar{X}_2$ has mean $\mu_1 - \mu_2$ and variance $\sigma_1^2/n_1 + \sigma_2^2/n_2$.

Proof. Linearity gives $E[\bar{X}] = n^{-1} \sum_i E[X_i] = \mu$; independence gives $\text{Var}(\bar{X}) = n^{-2} \sum_i \text{Var}(X_i) = \sigma^2/n$. For the difference, $\text{Var}(\bar{X}_1 - \bar{X}_2) = \text{Var}(\bar{X}_1) + \text{Var}(\bar{X}_2)$ by independence. $\square$

If the data are normal, $\bar{X} \sim \mathcal{N}(\mu, \sigma^2/n)$ exactly; in general the central limit theorem gives $\sqrt{n}(\bar{X} - \mu)/\sigma \xrightarrow{d} \mathcal{N}(0, 1)$.

Theorem (Sample proportion). If $Y \sim \text{Binomial}(n, p)$ and $\hat{p} = Y/n$, then $E[\hat{p}] = p$, $\text{Var}(\hat{p}) = p(1 - p)/n$, and $\sqrt{n}(\hat{p} - p)/\sqrt{p(1 - p)} \xrightarrow{d} \mathcal{N}(0, 1)$.

Proof. The moments follow from those of the binomial and the scaling rule $\text{Var}(Y/n) = \text{Var}(Y)/n^2$. The limit is the central limit theorem applied to $Y = \sum_i X_i$ with Bernoulli summands (the de Moivre–Laplace theorem). $\square$

Example (Classifier accuracy). A classifier with true accuracy $p = 0.8$ evaluated on $n = 100$ independent test examples has $\text{se}(\hat{p}) = \sqrt{0.8 \cdot 0.2/100} = 0.04$, so the observed accuracy lies in 0.8 ± 0.08 with probability about 95%.

4.2 Normal-theory distributions

Definition (Gaussian family). $X \sim \mathcal{N}(\mu, \sigma^2)$ has density $(2\pi\sigma^2)^{-1/2} \exp(-(x - \mu)^2/(2\sigma^2))$. Let $Z_1, \dots, Z_k \sim \mathcal{N}(0, 1)$ be independent. Then

- $\sum_i Z_i^2 \sim \chi_k^2$ (chi-square, k degrees of freedom), with mean k and variance $2k$;
- if $Z \sim \mathcal{N}(0, 1)$ and $V \sim \chi_k^2$ are independent, $Z/\sqrt{V/k} \sim t_k$ (Student);
- if $V_1 \sim \chi_a^2$ and $V_2 \sim \chi_b^2$ are independent, $(V_1/a)/(V_2/b) \sim F_{a,b}$.

Facts used below: a linear combination of jointly Gaussian variables is Gaussian; for jointly Gaussian variables, uncorrelated implies independent; $t_k \rightarrow \mathcal{N}(0, 1)$ as $k \rightarrow \infty$; $t_k^2 \sim F_{1,k}$.

Let $X_1, \dots, X_n \sim \mathcal{N}(\mu, \sigma^2)$ be i.i.d., with $\bar{X} = \bar{X}_n$ and $S^2 = (n-1)^{-1} \sum_i (X_i - \bar{X})^2$.

Theorem (Distribution of $\bar{X}$ and S^2). (i) $\bar{X} \sim \mathcal{N}(\mu, \sigma^2/n)$. (ii) $\bar{X}$ and S^2 are independent. (iii) $(n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$. (iv) Consequently

$$T = \frac{\bar{X} - \mu}{S/\sqrt{n}} \sim t_{n-1}. \quad (9)$$

Proof. (i) is closure of the Gaussian family under linear maps. For (ii)–(iii), let $Z_i = (X_i - \mu)/\sigma$ and Q an orthogonal $n \times n$ matrix whose first row is $(1/\sqrt{n}, \dots, 1/\sqrt{n})$. Then $Y = QZ$ has $Y_1 = \sqrt{n}\bar{Z}$ and, since Q is orthogonal, $Y \sim \mathcal{N}(0, I_n)$: the Y_i are independent standard normals. Orthogonality preserves the norm: $\sum_i Z_i^2 = \sum_i Y_i^2$, and $\sum_i Z_i^2 = \sum_i (Z_i - \bar{Z})^2 + n\bar{Z}^2$, hence

$$\sum_i (Z_i - \bar{Z})^2 = \sum_{i=2}^n Y_i^2 \sim \chi_{n-1}^2, \quad (10)$$

a function of $(Y_2, \dots, Y_n)$ only, and so independent of Y_1 and thus of $\bar{Z}$. Since $\sum_i (Z_i - \bar{Z})^2 = (n-1)S^2/\sigma^2$ and $\bar{X} = \mu + \sigma\bar{Z}$, (ii) and (iii) follow. For (iv), write $T = [(\bar{X} - \mu)/(\sigma/\sqrt{n})] / \sqrt{((n-1)S^2/\sigma^2)/(n-1)}$, a standard normal over the root of an independent $\chi_{n-1}^2/(n-1)$. $\square$

Theorem (Unbiasedness of S^2). For any i.i.d. sample with finite variance σ^2 , $\mathbb{E}[S^2] = \sigma^2$.

Proof. $\sum_i (X_i - \bar{X})^2 = \sum_i (X_i - \mu)^2 - n(\bar{X} - \mu)^2$. Taking expectations, $\mathbb{E}[\sum_i (X_i - \bar{X})^2] = n\sigma^2 - n(\sigma^2/n) = (n-1)\sigma^2$. $\square$

4.3 The t and F distributions in more detail

The t_k distribution has density $\Gamma((k+1)/2)/(\sqrt{k\pi}\Gamma(k/2)) \cdot (1 + t^2/k)^{-(k+1)/2}$. It is symmetric about 0, has mean 0 for $k > 1$ and variance $k/(k-2)$ for $k > 2$, and has heavier tails than $\mathcal{N}(0, 1)$, to which it converges as $k \rightarrow \infty$. Selected 97.5% quantiles are $t_{5,0.975} = 2.571$, $t_{9,0.975} = 2.262$, $t_{29,0.975} = 2.045$, against $z_{0.975} = 1.960$.

Lemma (Relation between t and F). If $T \sim t_k$ then $T^2 \sim F_{1,k}$.

Proof. Write $T = Z/\sqrt{V/k}$ with $Z \sim \mathcal{N}(0, 1)$ independent of $V \sim \chi_k^2$. Then $T^2 = Z^2/(V/k) = (Z^2/1)/(V/k)$ with $Z^2 \sim \chi_1^2$, the ratio defining $F_{1,k}$. $\square$

The $F_{a,b}$ distribution has mean $b/(b-2)$ for $b > 2$ and is the null distribution of the ratio of two independent normal-sample variance estimates, S_1^2/S_2^2 scaled by σ_2^2/σ_1^2 , with $a = n_1 - 1$ and $b = n_2 - 1$; it is also the null distribution of the overall test in analysis of variance and linear regression.

4.4 A non-normal example: the sample maximum

Theorem (Maximum of uniform samples). Let $X_1, \dots, X_n$ be i.i.d. $\text{Uniform}(0, \theta)$ and $M = \max_i X_i$. Then $P(M \leq m) = (m/\theta)^n$ for $0 \leq m \leq \theta$, with density nm^{n-1}/θ^n , $E[M] = n\theta/(n+1)$, and $(n+1)M/n$ is unbiased for θ .

Proof. $M \leq m$ iff every $X_i \leq m$, so by independence $P(M \leq m) = (m/\theta)^n$. Differentiating gives the density, and $E[M] = \int_0^\theta m \cdot nm^{n-1}/\theta^n dm = n\theta/(n+1)$. $\square$

The sampling distribution of M is concentrated against the boundary θ and is far from normal, so normal-theory intervals do not apply: sampling distributions must be derived for the statistic actually used.

4.5 Simulation and the bootstrap

Any sampling distribution can be approximated by simulation: draw B samples from the assumed model, compute T on each, and use the empirical distribution of $T_1, \dots, T_B$; by the law of large numbers its mean and standard deviation converge to $E[T]$ and $\text{se}(T)$ as $B \rightarrow \infty$.

When the model is unknown, the bootstrap replaces the unknown distribution F by the empirical distribution $\hat{F}_n$, which puts mass $1/n$ on each observation. Resample n observations with replacement, compute T_b^* , repeat for $b = 1, \dots, B$, and estimate $\text{se}(T)$ by the standard deviation of the T_b^* . It is consistent for smooth statistics such as the mean; it fails for extreme statistics such as the maximum in the previous example.

5 Point estimation

This part uses the probability toolkit (axioms, expectation and variance rules, limit theorems, Slutsky) and the sampling distributions of the mean and variance, the χ^2 , t and F laws, from the chapters above. They are quoted without proof.

Let $X_1, \dots, X_n \sim f(\cdot; \theta)$, $\theta \in \Theta \subset \mathbb{R}$, and $\hat{\theta} = T(X_1, \dots, X_n)$.

Definition. $\text{bias}(\hat{\theta}) = E_\theta[\hat{\theta}] - \theta$. The estimator is unbiased if this vanishes for all θ . It is consistent if $\hat{\theta}_n \xrightarrow{P} \theta$. Its mean squared error is $\text{MSE}(\hat{\theta}) = E_\theta[(\hat{\theta} - \theta)^2]$.

Theorem (Bias-variance decomposition).

$$\text{MSE}(\hat{\theta}) = \text{Var}_\theta(\hat{\theta}) + \text{bias}(\hat{\theta})^2. \quad (11)$$

Proof. Add and subtract $E\hat{\theta}$: $(\hat{\theta} - \theta)^2 = (\hat{\theta} - E\hat{\theta})^2 + 2(\hat{\theta} - E\hat{\theta})(E\hat{\theta} - \theta) + (E\hat{\theta} - \theta)^2$. The cross term has zero expectation because $E\hat{\theta} - \theta$ is a constant and $E[\hat{\theta} - E\hat{\theta}] = 0$. $\square$

Example (A biased estimator with smaller MSE). For $X_i \sim \mathcal{N}(\mu, \sigma^2)$, the estimator $\hat{\sigma}_c^2 = c \sum_i (X_i - \bar{X})^2$ has, using $\sum_i (X_i - \bar{X})^2 \sim \sigma^2 \chi_{n-1}^2$ (mean $(n-1)\sigma^2$, variance $2(n-1)\sigma^4$),

$$\text{MSE} = \sigma^4[2c^2(n-1) + ((n-1)c-1)^2]. \quad (12)$$

This is minimized at $c = 1/(n + 1)$, not at the unbiased choice $c = 1/(n - 1)$. The unbiased estimator is therefore not MSE-optimal, a first instance of the shrinkage idea behind Ridge and Lasso.

5.1 Fisher information and the Cramér–Rao bound

Assume the regularity conditions: the support of $f(\cdot; \theta)$ does not depend on θ , $\log f$ is twice differentiable in θ , and differentiation under the integral sign is allowed.

Definition (Score and Fisher information). The score of one observation is $s(x; \theta) = \partial_\theta \log f(x; \theta)$, and

$$I(\theta) = E_\theta[s(X; \theta)^2] = \text{Var}_\theta(s(X; \theta)). \quad (13)$$

Lemma. Under the regularity conditions, $E_\theta[s(X; \theta)] = 0$ and $I(\theta) = -E_\theta[\partial_\theta^2 \log f(X; \theta)]$. For n independent observations the information is $nI(\theta)$.

Proof. $E[s] = \int (\partial_\theta f / f) f dx = \partial_\theta \int f dx = \partial_\theta 1 = 0$. Differentiating $\int s f dx = 0$ once more gives $\int (\partial_\theta s) f dx + \int s^2 f dx = 0$, and $\partial_\theta s = \partial_\theta^2 \log f$. Additivity follows because the log-likelihood of independent data is the sum of the individual ones, so the scores add and, being independent with mean zero, their variances add. $\square$

Theorem (Cramér–Rao lower bound). If $\hat{\theta}$ is unbiased with finite variance, then $\text{Var}_\theta(\hat{\theta}) \geq 1/(nI(\theta))$.

Proof. Let $U = \sum_i s(X_i; \theta)$ be the total score, so $EU = 0$ and $\text{Var}(U) = nI(\theta)$. Unbiasedness reads $\int \hat{\theta}(x) f_{n(x; \theta)} dx = \theta$ (f_n the joint density). Differentiating in θ gives $\int \hat{\theta}(x) U(x) f_n dx = 1$, i.e. $\text{Cov}(\hat{\theta}, U) = E[\hat{\theta}U] = 1$ since $EU = 0$. Cauchy–Schwarz: $1 = \text{Cov}(\hat{\theta}, U)^2 \leq \text{Var}(\hat{\theta}) \text{Var}(U) = \text{Var}(\hat{\theta}) nI(\theta)$. $\square$

Remark. For a biased estimator with $b(\theta) = \text{bias}(\hat{\theta})$ the same argument gives $\text{Var}(\hat{\theta}) \geq (1 + b'(\theta))^2 / (nI(\theta))$. Equality in Cauchy–Schwarz holds iff $\hat{\theta} - \theta$ is proportional to the score, which happens exactly in exponential families for the mean-value parameter.

5.2 Sufficiency (brief)

A statistic $T(X)$ is sufficient for θ if the conditional law of the sample given T does not depend on θ . By the Fisher–Neyman factorization theorem, T is sufficient iff $f_{n(x; \theta)} = g(T(x); \theta)h(x)$. For normal data with known σ , $\sum_i X_i$ is sufficient for μ ; for the Bernoulli model, $\sum_i X_i$ is sufficient for p . The Rao–Blackwell theorem states that conditioning an estimator on a sufficient statistic never increases its MSE.

6 Maximum likelihood

Definition. The likelihood is $L(\theta) = \prod_{i=1}^n f(x_i; \theta)$, the log-likelihood $\ell(\theta) = \log L(\theta)$ and the maximum likelihood estimator $\hat{\theta}_{\text{MLE}} = \arg \max_\theta \ell(\theta)$.

Example (Bernoulli). $\ell(p) = (\sum x_i) \log p + (n - \sum x_i) \log(1 - p)$, $\ell'(p) = (\sum x_i)/p - (n - \sum x_i)/(1 - p)$. Setting it to zero gives $\hat{p} = \bar{x}$; $\ell'' < 0$ so it is a maximum. The information is $I(p) = 1/(p(1 - p))$ and $\text{Var}(\hat{p}) = p(1 - p)/n = 1/(nI(p))$: the MLE attains the Cramér–Rao bound exactly.

Example (Normal). $\ell(\mu, \sigma^2) = -(n/2) \log(2\pi\sigma^2) - (2\sigma^2)^{-1} \sum_i (x_i - \mu)^2$. The first-order conditions give $\hat{\mu} = \bar{x}$ and $\hat{\sigma}^2 = n^{-1} \sum_i (x_i - \bar{x})^2$. By the theorem on S^2 , $E\hat{\sigma}^2 = ((n - 1)/n)\sigma^2$: the MLE of the variance is biased, but consistent, and $\hat{\sigma}^2 = cS^2$ with $c = (n - 1)/n \rightarrow 1$.

6.1 Asymptotic properties

Theorem (Asymptotics of the MLE). Under the regularity conditions above, with θ_0 the true parameter and $I(\theta_0) > 0$: (i) $\hat{\theta}_{\text{MLE}} \xrightarrow{P} \theta_0$; (ii) $\sqrt{n}(\hat{\theta}_{\text{MLE}} - \theta_0) \xrightarrow{d} \mathcal{N}(0, I(\theta_0)^{-1})$; (iii) invariance: the MLE of $g(\theta)$ is $g(\hat{\theta}_{\text{MLE}})$, and by the delta method it is asymptotically $\mathcal{N}(g(\theta_0), g'(\theta_0)^2/(nI(\theta_0)))$.

Proof. Sketch of (ii). Expand the score at the MLE around θ_0 : $0 = \ell'(\hat{\theta}) = \ell'(\theta_0) + (\hat{\theta} - \theta_0)\ell''(\tilde{\theta})$ for some $\tilde{\theta}$ between them. Hence

$$\sqrt{n}(\hat{\theta} - \theta_0) = \frac{n^{-1/2}\ell'(\theta_0)}{-n^{-1}\ell''(\tilde{\theta})}. \quad (14)$$

The numerator is $n^{-1/2} \sum_i s(X_i; \theta_0)$, a normalized sum of i.i.d. mean-zero variables with variance $I(\theta_0)$, so by the CLT it tends to $\mathcal{N}(0, I(\theta_0))$. By the law of large numbers and consistency, the denominator tends to $-E[\partial_{\theta}^2 \log f] = I(\theta_0)$. Slutsky gives the limit $\mathcal{N}(0, I(\theta_0)^{-1})$. $\square$

The practical consequence is the standard error $\hat{s}e = 1/\sqrt{nI(\hat{\theta})}$ (or its observed-information version, $1/\sqrt{-\ell''(\hat{\theta})}$), valid for large n .

6.2 Method of moments

Equate the first d population moments $m_{j(\theta)} = E_{\theta} X^j$ to the sample moments $n^{-1} \sum_i X_i^j$ and solve for $\theta \in \mathbb{R}^d$. The estimator is consistent when $m = (m_1, \dots, m_d)$ is invertible with continuous inverse, but generally has larger asymptotic variance than the MLE.

7 Interval estimation

Definition. A $(1 - \alpha)$ confidence interval is a random interval $[L(X), U(X)]$ with $P_{\theta}(L(X) \leq \theta \leq U(X)) \geq 1 - \alpha$ for all $\theta \in \Theta$. The left side is the coverage; the inequality is replaced by an equality for exact intervals.

The probability refers to the random endpoints under repeated sampling. For the realized interval, θ is fixed and the event either holds or not.

Definition (Pivot). $Q(X; \theta)$ is a pivot if its distribution does not depend on θ .

If $P(a \leq Q(X; \theta) \leq b) = 1 - \alpha$ and Q is monotone in θ , solving $a \leq Q \leq b$ for θ gives an exact interval.

Theorem (Intervals for the normal model). Let $X_i \sim \mathcal{N}(\mu, \sigma^2)$ i.i.d.. Exact $(1 - \alpha)$ intervals:

- mean, σ unknown: $\bar{X} \pm t_{n-1, 1-\alpha/2} S / \sqrt{n}$;
- mean, σ known: $\bar{X} \pm z_{1-\alpha/2} \sigma / \sqrt{n}$;
 $[(n-1)S^2 / \chi_{n-1, 1-\alpha/2}^2, (n-1)S^2 / \chi_{n-1, \alpha/2}^2]$
- variance: $(n-1)S^2 / \chi_{n-1, \alpha/2}^2$

Proof. $P(-t \leq T \leq t) = 1 - \alpha$ with $t = t_{n-1, 1-\alpha/2}$ and T the pivot of the sampling-distribution theorem; rearranging $-t \leq (\bar{X} - \mu)/(S/\sqrt{n}) \leq t$ gives the first interval. The second is identical with T replaced by $\sqrt{n}(\bar{X} - \mu)/\sigma \sim \mathcal{N}(0, 1)$. For the third, $P(\chi_{n-1, \alpha/2}^2 \leq (n-1)S^2/\sigma^2 \leq \chi_{n-1, 1-\alpha/2}^2) = 1 - \alpha$ and inverting in σ^2 reverses the inequalities. $\square$

Theorem (Wald interval). If $\sqrt{n}(\hat{\theta} - \theta) \xrightarrow{d} \mathcal{N}(0, v(\theta))$ with v continuous, then $\hat{\theta} \pm z_{1-\alpha/2} \sqrt{v(\hat{\theta})/n}$ has asymptotic coverage $1 - \alpha$. For the MLE, $v = I^{-1}$. For a proportion, $\hat{p} \pm z_{1-\alpha/2} \sqrt{\hat{p}(1-\hat{p})/n}$.

Proof. $\sqrt{n}(\hat{\theta} - \theta)/\sqrt{v(\hat{\theta})} = [\sqrt{n}(\hat{\theta} - \theta)/\sqrt{v(\theta)}] \cdot [\sqrt{v(\theta)/v(\hat{\theta})}]$; the first factor tends in distribution to $\mathcal{N}(0, 1)$ and the second in probability to 1 by consistency and continuity. Apply Slutsky. $\square$

Remark. The Wald interval for a proportion has coverage that oscillates and can be far below $1 - \alpha$ when n is small or p is near 0 or 1. The Wilson interval, obtained by inverting the score test instead of replacing p by $\hat{p}$ in the variance, is the standard fix.

8 Hypothesis testing

8.1 Framework

A test of $H_0 : \theta \in \Theta_0$ against $H_1 : \theta \in \Theta_1$ is a rule $\varphi(X) \in \{0, 1\}$ (1 means reject), typically $\varphi = \mathbb{1}\{T(X) \in R\}$ for a statistic T and rejection region R .

Definition. The size of the test is $\sup_{\theta \in \Theta_0} P_\theta(\varphi = 1)$; it has level α if the size is at most α . The power function is $\beta(\theta) = P_\theta(\varphi = 1)$, so the Type II error at $\theta \in \Theta_1$ is $1 - \beta(\theta)$.

Definition. The p -value of observed t_{obs} for a statistic T that rejects for large values is $p = \sup_{\theta \in \Theta_0} P_\theta(T \geq t_{\text{obs}})$.

Lemma (Uniformity of the p -value). If H_0 is simple, T has a continuous distribution function G under H_0 , and rejection is for large T , then $p(T) = 1 - G(T) \sim \text{Uniform}(0, 1)$ under H_0 . Hence rejecting when $p \leq \alpha$ has size exactly α .

Proof. $G(T) \sim \text{Uniform}(0, 1)$ by the probability integral transform: $P(G(T) \leq u) = P(T \leq G^{-1}(u)) = u$. Then $1 - G(T)$ is also uniform. $\square$

8.2 The t -tests

Theorem (One-sample and Welch t -tests). (i) For $X_i \sim \mathcal{N}(\mu, \sigma^2)$ i.i.d., under $H_0 : \mu = \mu_0$ the statistic $T = (\bar{X} - \mu_0)/(S/\sqrt{n})$ has the t_{n-1} distribution, so rejecting for $|T| > t_{n-1, 1-\alpha/2}$ has size α . (ii) For two independent normal samples with means μ_1, μ_2 , and $H_0 : \mu_1 = \mu_2$,

$$T = \frac{\bar{X}_1 - \bar{X}_2}{\sqrt{S_1^2/n_1 + S_2^2/n_2}} \quad (15)$$

is approximately t_ν with

$$\nu = \frac{(S_1^2/n_1 + S_2^2/n_2)^2}{(S_1^2/n_1)^2/(n_1 - 1) + (S_2^2/n_2)^2/(n_2 - 1)}. \quad (16)$$

Part (i) is the pivot theorem with $\mu = \mu_0$. In (ii) the denominator is a weighted sum of independent scaled χ^2 variables, which is not exactly a scaled χ^2 ; the Welch-Satterthwaite ν matches the first two moments of a scaled χ_ν^2 to it. With equal variances the pooled test is exactly $t_{n_1+n_2-2}$, but Welch is robust to unequal variances at little cost.

Theorem (Duality of tests and intervals). If, for each θ_0 , $A(\theta_0)$ is the acceptance region of a level- α test of $\theta = \theta_0$, then $C(x) = \{\theta_0 : x \in A(\theta_0)\}$ is a $(1 - \alpha)$ confidence set. Conversely, from a confidence set one obtains a test by rejecting $\theta = \theta_0$ iff $\theta_0 \notin C(x)$.

Proof. $P_{\theta_0}(\theta_0 \in C(X)) = P_{\theta_0}(X \in A(\theta_0)) \geq 1 - \alpha$, for every θ_0 . The converse is the same identity read backwards. $\square$

8.3 Power and sample size

Theorem (Power of the one-sided z -test). Let $X_i \sim \mathcal{N}(\mu, \sigma^2)$ with σ known, and reject $H_0 : \mu = \mu_0$ in favour of $H_1 : \mu > \mu_0$ when $Z = \sqrt{n}(\bar{X} - \mu_0)/\sigma > z_{1-\alpha}$. Then

$$\beta(\mu_1) = 1 - \Phi(z_{1-\alpha} - \sqrt{n}(\mu_1 - \mu_0)/\sigma), \quad (17)$$

and $\beta(\mu_1) \geq 1 - \beta_0$ holds iff $n \geq \left((z_{1-\alpha} + z_{1-\beta_0})\sigma/(\mu_1 - \mu_0) \right)^2$.

Proof. Under $\mu = \mu_1$, $\sqrt{n}(\bar{X} - \mu_1)/\sigma \sim \mathcal{N}(0, 1)$, so $Z = Z_0 + \sqrt{n}(\mu_1 - \mu_0)/\sigma$ with $Z_0 \sim \mathcal{N}(0, 1)$. Hence $P(Z > z_{1-\alpha}) = 1 - \Phi(z_{1-\alpha} - \sqrt{n}(\mu_1 - \mu_0)/\sigma)$. Requiring this to be at least $1 - \beta_0$ means $z_{1-\alpha} - \sqrt{n}\delta/\sigma \leq -z_{1-\beta_0}$ (with $\delta = \mu_1 - \mu_0$), and solving for n gives the bound. $\square$

The required n scales like σ^2/δ^2 : halving the detectable effect quadruples the sample.

8.4 Likelihood ratio tests

Theorem (Neyman–Pearson lemma). For simple $H_0 : \theta = \theta_0$ against simple $H_1 : \theta = \theta_1$, let φ^* reject when $L(\theta_1)/L(\theta_0) > k$, with k chosen so that φ^* has size α . Then for any other test φ of size at most α , $\beta_\varphi(\theta_1) \leq \beta_{\varphi^*}(\theta_1)$.

Proof. For every x , $(\varphi^*(x) - \varphi(x))(f(x; \theta_1) - kf(x; \theta_0)) \geq 0$: where $\varphi^* = 1$ the second factor is positive, where $\varphi^* = 0$ it is non-positive and the first factor is $-\varphi(x) \leq 0$ or zero. Integrating over x gives $\beta_{\varphi^*}(\theta_1) - \beta_\varphi(\theta_1) \geq k(\alpha_{\varphi^*} - \alpha_\varphi) \geq 0$, since both sizes are at most α and φ^* has size exactly α . $\square$

Theorem (Wilks' theorem). Let $\Lambda = \sup_{\Theta_0} L / \sup_{\Theta} L$, where Θ has d free parameters and Θ_0 has $d - r$. Under the regularity conditions and H_0 ,

$$-2 \log \Lambda \xrightarrow{d} \chi_r^2. \quad (18)$$

The proof expands ℓ quadratically around the MLE and uses asymptotic normality of the MLE, so that $-2 \log \Lambda$ becomes a quadratic form in r asymptotically independent standard normals. The Wald and score tests are asymptotically equivalent to it. In the linear model with Gaussian errors, the likelihood ratio is a monotone function of an F statistic, which is why the F -test of regression has an exact distribution.

8.5 Multiple testing

With m tests and all nulls true, each at level α , the expected number of false rejections is $m\alpha$.

Theorem (Bonferroni and Benjamini–Hochberg). (i) Rejecting each hypothesis with $p_i \leq \alpha/m$ controls the family-wise error rate, $P(\text{at least one false rejection}) \leq \alpha$, with no assumption on the dependence of the p -values. (ii) Order the p -values $p_{(1)} \leq \dots \leq p_{(m)}$ and reject the k smallest, where $k = \max\{i : p_{(i)} \leq i\alpha/m\}$. If the p -values are independent, the false discovery rate $E[V/\max(R, 1)]$ (with V false rejections out of R) is at most $\alpha m_0/m \leq \alpha$, where m_0 is the number of true nulls.

Proof. (i) By the union bound, $P\left(\bigcup_{i \in H_0} \{p_i \leq \alpha/m\}\right) \leq m_0\alpha/m \leq \alpha$. (ii) is the Benjamini–Hochberg (1995) theorem; the proof conditions on the other p -values and is omitted. $\square$

9 Summary of the main formulas

Quantity	Result
Standard error of the mean	$\sigma/\sqrt{n}$
S^2 unbiased	$E S^2 = \sigma^2$, divisor $n - 1$

Normal sampling laws	$(n-1)S^2/\sigma^2 \sim \chi_{n-1}^2, T \sim t_{n-1}, \bar{X} \perp S^2$
MSE	$\text{Var} + \text{bias}^2$
Cramér–Rao	$\text{Var}(\hat{\theta}) \geq 1/(nI(\theta))$
MLE	$\sqrt{n}(\hat{\theta} - \theta_0) \xrightarrow{d} \mathcal{N}(0, I(\theta_0)^{-1})$
CI for the mean	$\bar{X} \pm t_{n-1, 1-\alpha/2} S/\sqrt{n}$
CI for the variance	$[(n-1)S^2/\chi_{n-1, 1-\alpha/2}^2, (n-1)S^2/\chi_{n-1, \alpha/2}^2]$
Wald interval	$\hat{\theta} \pm z_{1-\alpha/2} \hat{\text{se}}$
Sample size, one-sided z -test	$n = ((z_{1-\alpha} + z_{1-\beta})\sigma/\delta)^2$
Likelihood ratio	$-2 \log \Lambda \xrightarrow{d} \chi_r^2$
Bonferroni / BH	$p_i \leq \alpha/m$ / largest k with $p_{(k)} \leq k\alpha/m$

10 References

- G. Casella and R. L. Berger, *Statistical Inference*, 2nd ed., Duxbury, 2002: the standard source for the probability and distribution theory (chapters 1 to 5) and for the proofs of the Cramér–Rao bound, the MLE asymptotics and the Neyman–Pearson lemma.
- J. L. Devore, *Probability and Statistics for Engineering and the Sciences*, Cengage, 2011: the course text for the first part of the syllabus (chapters 1 to 6).
- L. Wasserman, *All of Statistics*, Springer, 2004: concise statements of the limit theorems, the bootstrap, MLE and multiple testing.
- B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*, Chapman and Hall, 1993.
- Y. Benjamini and Y. Hochberg, “Controlling the false discovery rate”, *J. R. Statist. Soc. B* 57 (1995).